




ICML

International Conference
On Machine Learning

MECAT: A Multi-Experts Constructed Benchmark for Fine-Grained Audio Understanding Tasks

Yadong Niu^{*1}, Tianzi Wang^{*1,2}, Heinrich Dinkel¹, Xingwei Sun¹, Jiahao Zhou¹, Gang Li¹, Jizhong Liu¹, Xunying Liu², Jian Luan¹

¹MiLM Plus, Xiaomi Inc ²The Chinese University of Hong Kong ^{*}Equal Contribution



Highlights

Motivation.

- Existing audio benchmarks fail to assess fine-grained understanding.
- Annotations are often coarse, generic, or single-perspective.
- Evaluation metrics penalize novel details while tolerating overly safe descriptions.


Solution.

- MECAT builds a multi-expert pipeline for detailed audio understanding.
- LLM-CoT synthesizes expert outputs into dense captions and QA pairs.
- DATE rewards specific details and penalizes generic terms.

Key advantages.

- Scale:** **20K** clips from ACAV100M across extended audio domains.
- Density (Per Clip):**
 - Caption: **18** refs spanning systemic, content-related/unrelated aspects.
 - QA: **5** pairs probing perception, analysis, and reasoning capabilities.
- Discrimination:** DATE robustly separates accurate details from vague errors.

Construction Framework & Taxonomy




Diversity & Quality Validation

Diversity

- Top Row (Audio): Broader feature space across 8 extended domains.
- Bottom Row (Text): Richer semantic clustering across multiple categories.

Quality

- Multi-stage filtering keeps only the top **10%** candidates.
- Expert audit confirms **3.81/5** relevance with only **3.4%** strict error rate.



Metric Analysis


DATE penalizes generic descriptions while rewarding fine-grained accuracy.

- Mechanism:** single-sample similarity and cross-sample discriminability.
- Alignment:** align with human preference (preferred: 90.9 vs. non-preferred: 49.3).

$$DATE_i = \frac{2 \cdot S_{sim,i} \cdot S_{dis,i}}{S_{sim,i} + S_{dis,i}} \in [0, 1]$$

Proof (Qualitative & Quantitative)

- A:** Traditional metrics **struggle to separate** “Safe” from “Right”.
- B:** DATE mimics LLM-as-judge to **widen the gap** between “Safe” and “Right”.



Benchmarking Results (Representative)

Captioning: Top LALMs improve overall scores, but

- exhibit a **severe speech-centric bias**
- struggle to describe content-unrelated acoustic environments.

Model	Systemic		Speech		Music		Sound		Env	Score _{Cap}
	Long	Short	Pure	Mixed	Pure	Mixed	Pure	Mixed		
Kimi-Audio	49.5	54.2	30.0	31.3	27.7	16.9	43.1	16.2	7.0	32.8
Qwen2.5-Omni	61.1	56.5	39.9	40.9	32.1	30.9	50.7	23.8	17.9	42.6
Qwen3-Omni-Flash	65.7	62.5	59.2	59.9	57.4	32.5	55.8	31.6	27.1	52.9
Gemini-3-Pro	64.9	65.8	60.5	62.4	49.8	39.8	51.6	30.7	26.1	53.1

QA: Models handle semantic questions well, but

- struggle with **acoustic quality** and **sound-characteristic reasoning**.
- fail to surpass the 40-point threshold in Quality Assessment.

Model	Direct Perception	Sound Characteristics	Quality Assessment	Inference Judgment	Application Context	Score _{QA}
Kimi-Audio	45.6	39.2	18.7	48.9	41.2	38.0
Qwen2.5-Omni	57.8	52.9	39.1	53.2	50.8	49.6
Qwen3-Omni-Flash	48.0	45.9	29.5	56.7	54.8	46.7
Gemini-3-Pro	55.5	54.4	37.7	57.3	56.6	51.5

Conclusion & Contact

Conclusion & Future work:

- Contributions:** Dense benchmark and DATE metric for fine-grained details.
- Insights:** Exposes LALM weakness in speech bias and mixed scenes.
- Limitations:** 10s clip length restricts long-range temporal evaluation.
- Future Work:** Advancing audio understanding and generation frameworks.

Contact & Connect



Scan for Data & Code



LinkedIn



WeChat